



universidade de aveiro

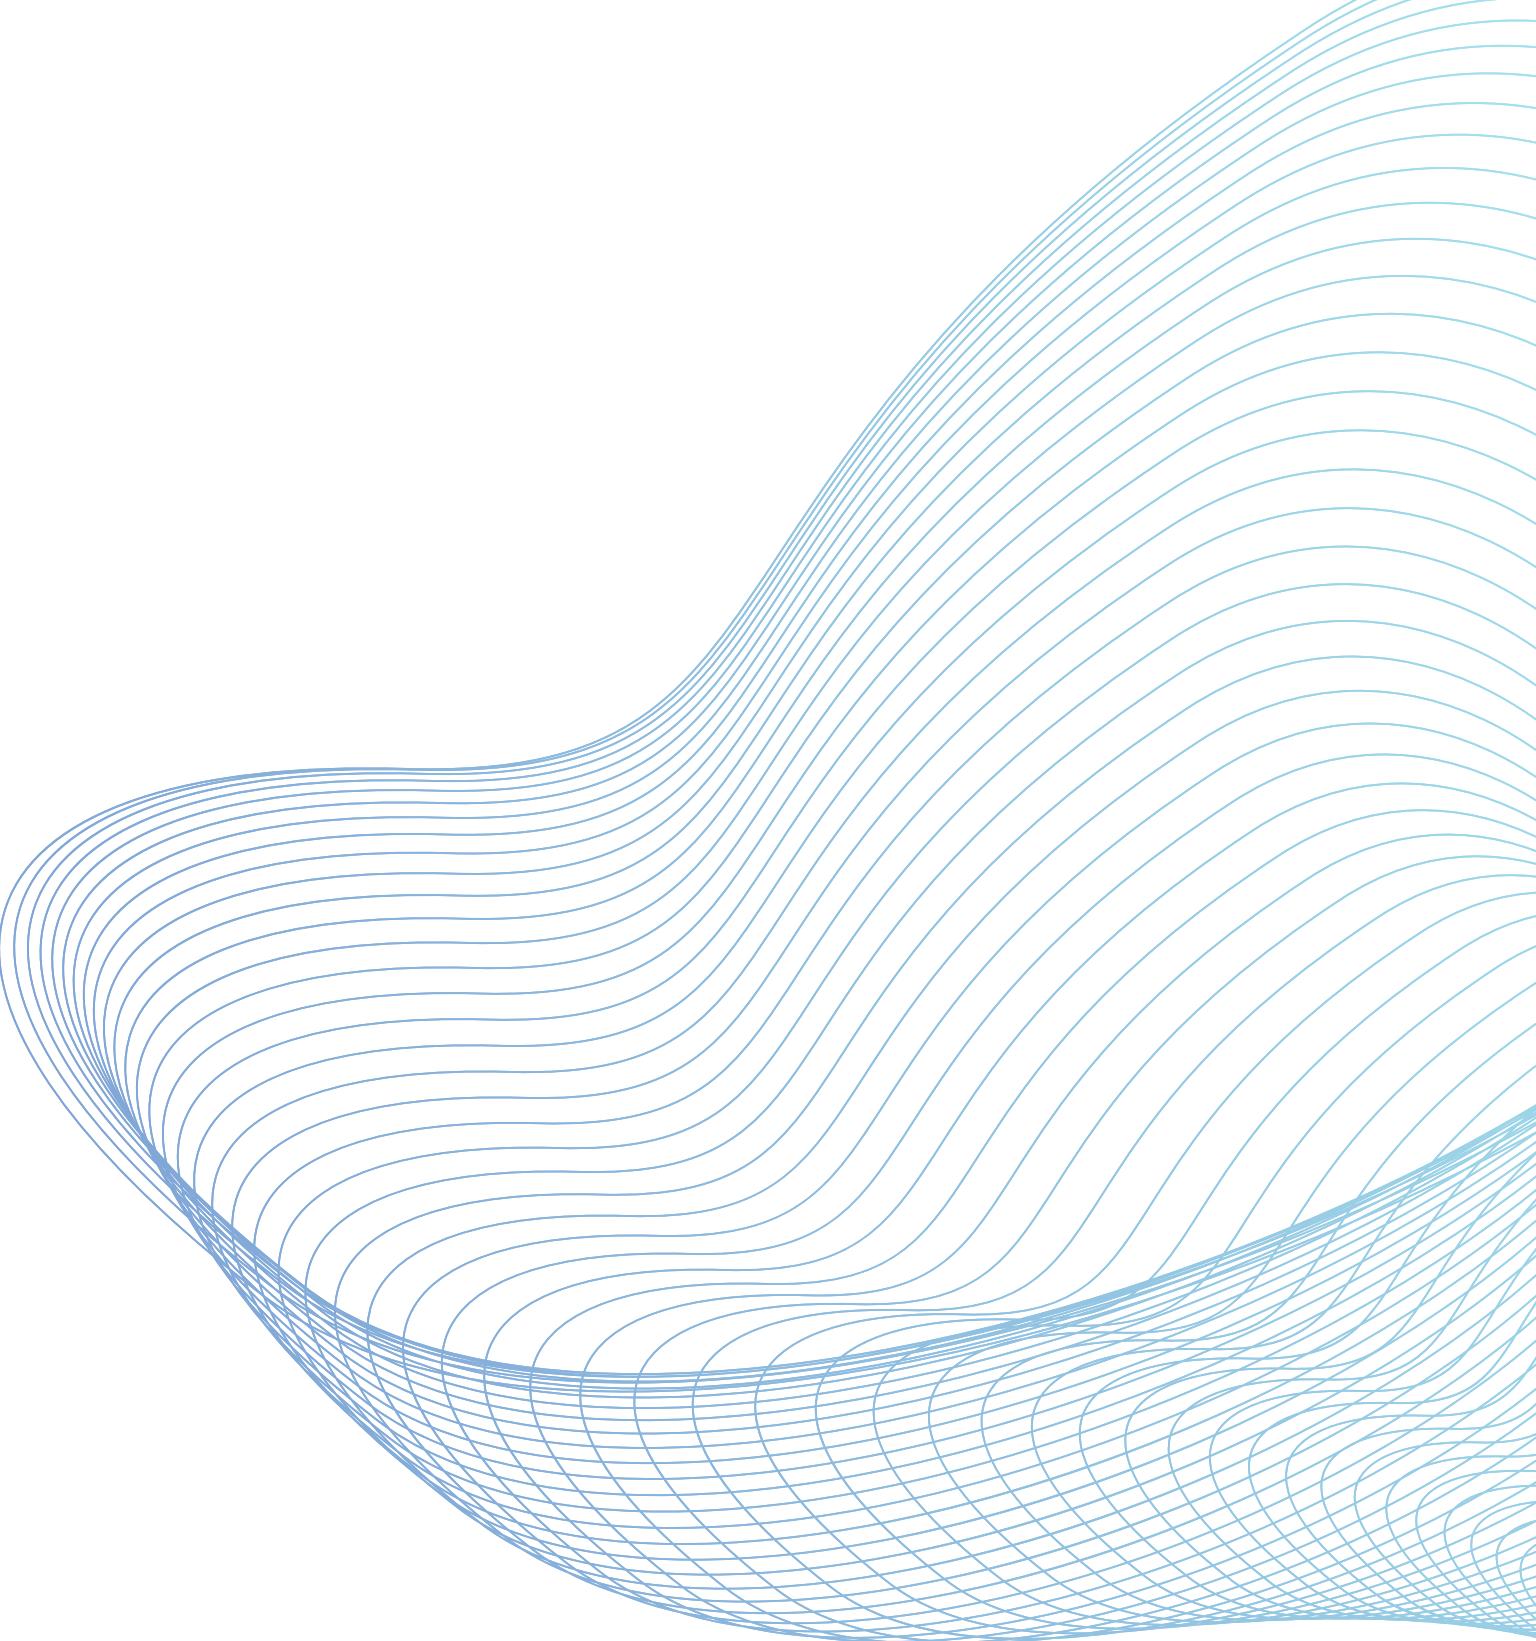
TENSORFLOW AGENTS

 - Agents

The TensorFlow Agents logo consists of a stylized orange 'T' icon followed by the word '- Agents' in a large, bold, yellow sans-serif font.

Daniel Emídio (108986)
David Palricas (108780)
Márcio Tavares (108096)

Professor
Luís Seabra Lopes
(lsl@ua.pt)



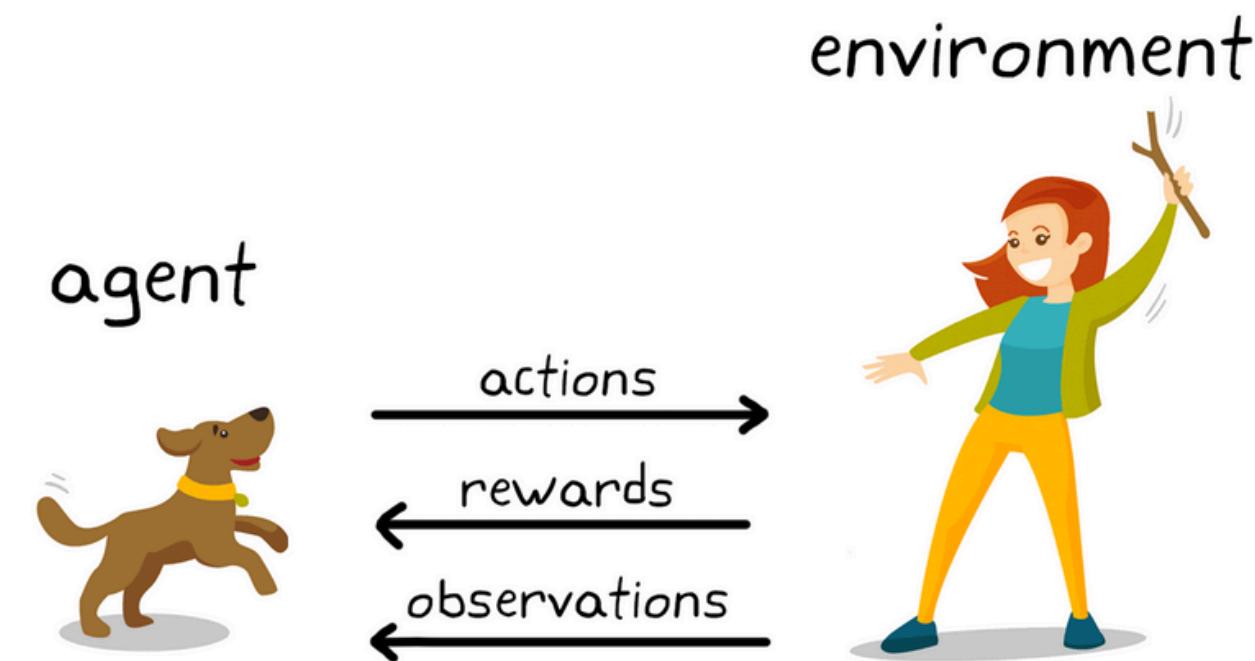
INTRODUCTION



Language:

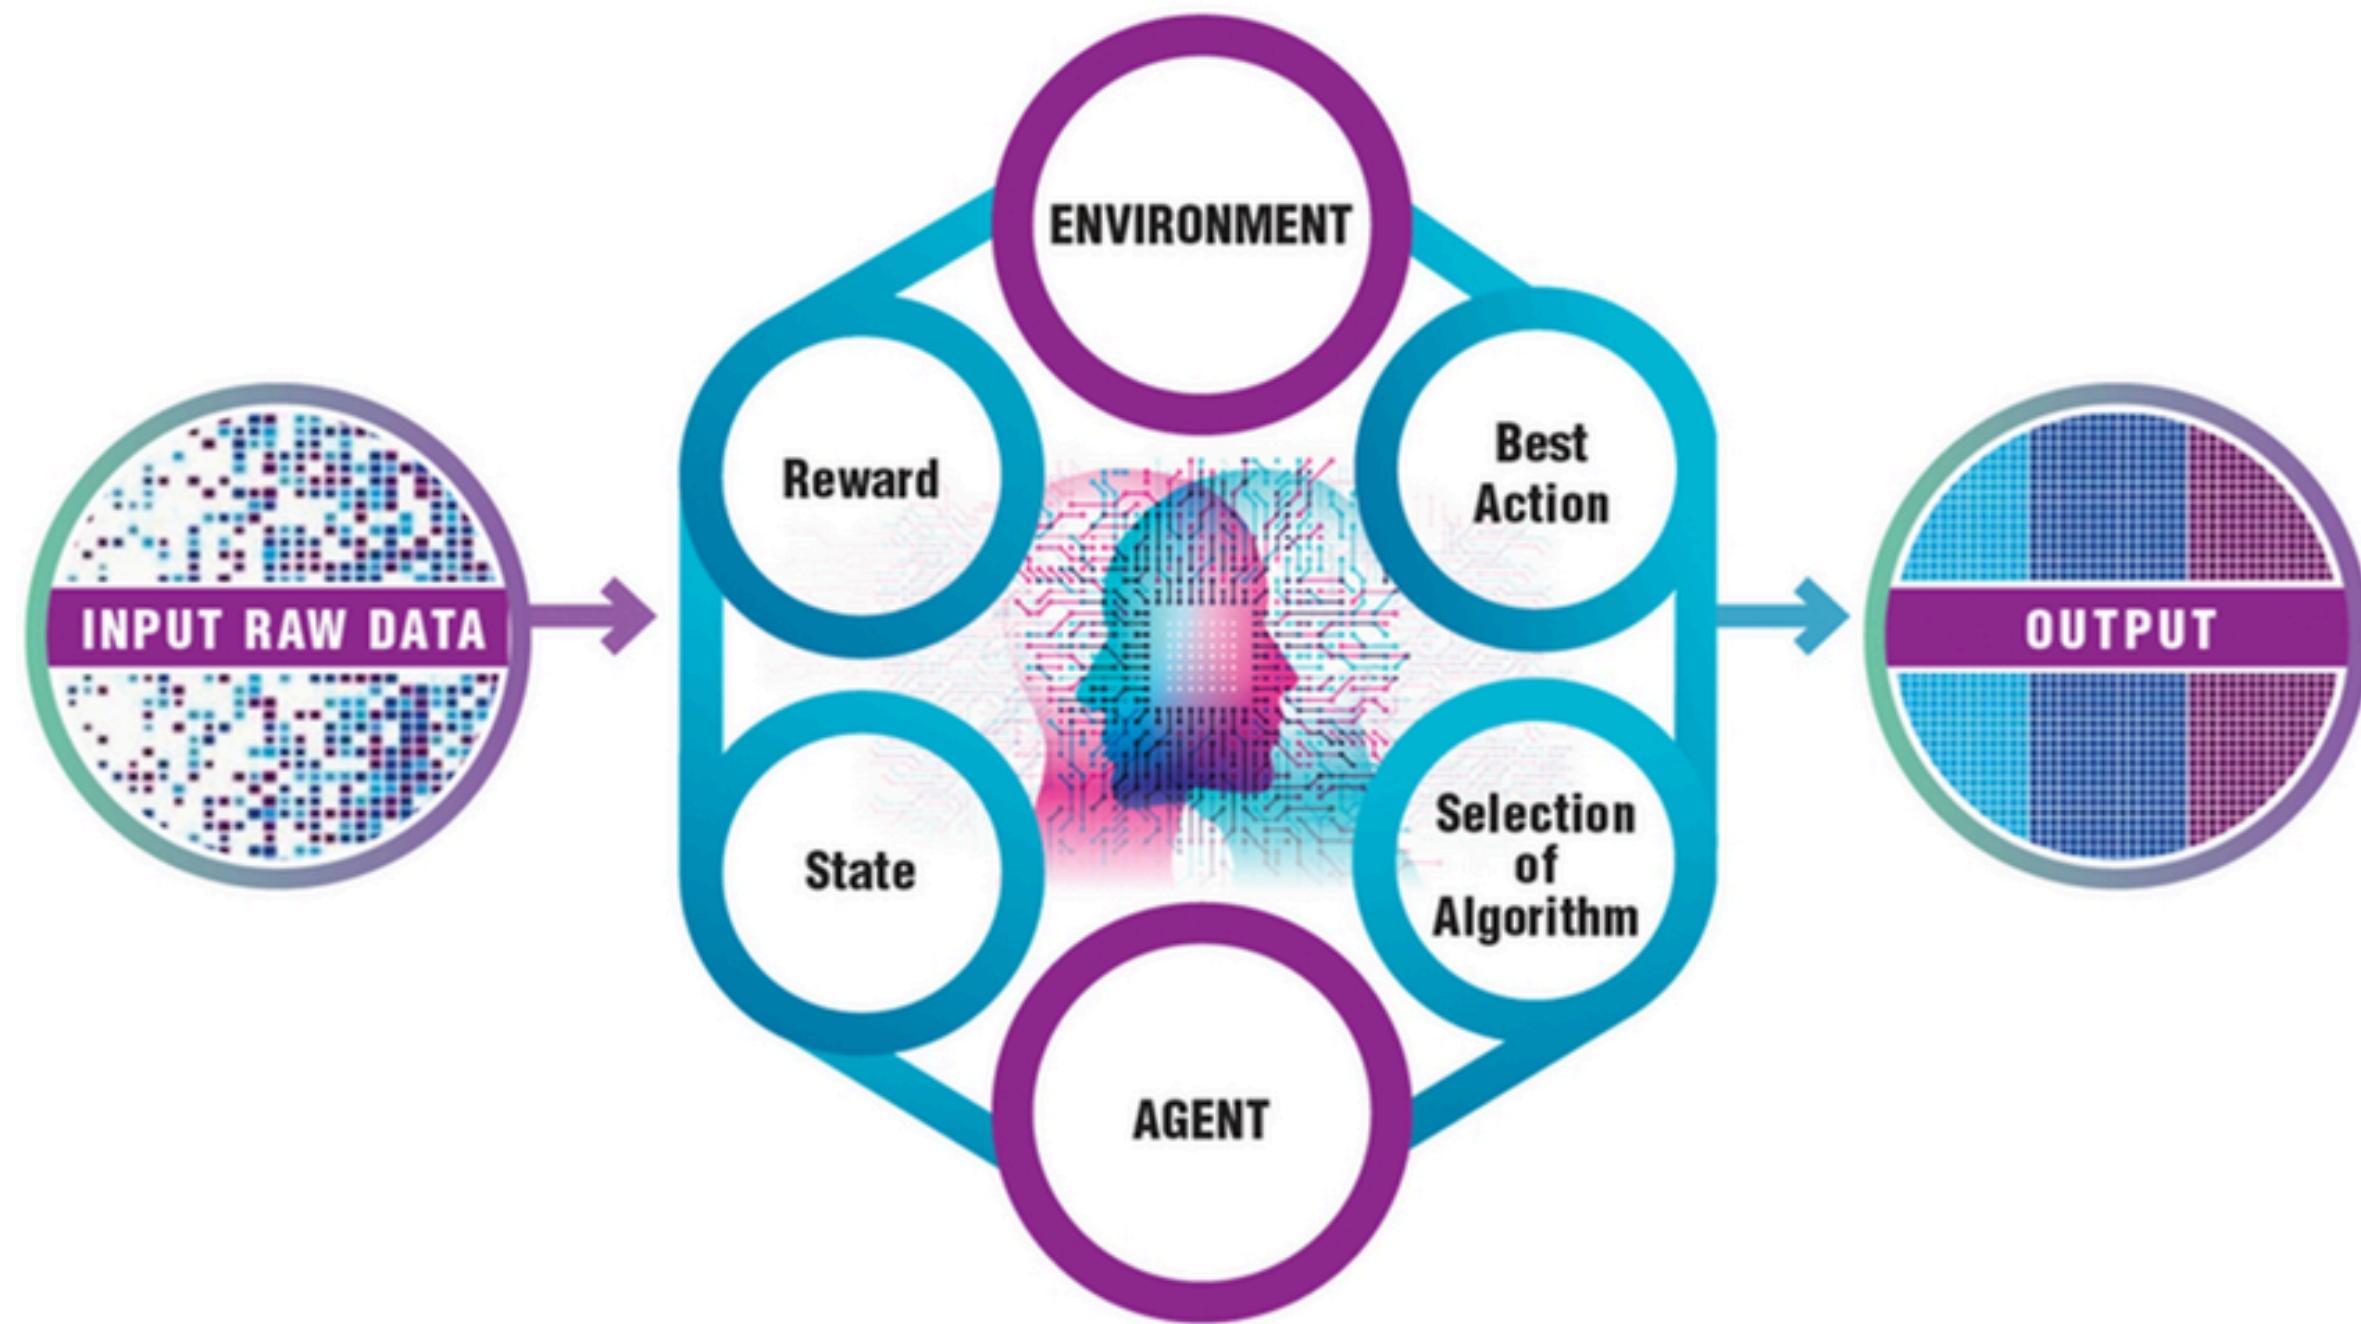


Reinforcement Learning:



BASE PRINCIPLES

1. Agent
2. Environment
3. State
4. Action
5. Reward
6. Policy



USE CASES

Used By:

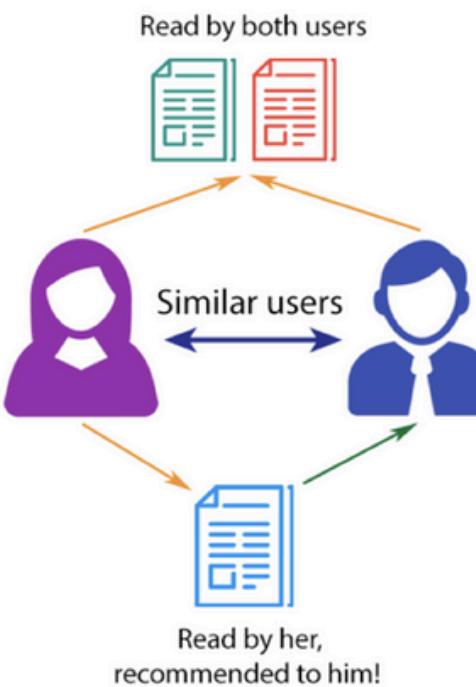
Google

Uber



Recommendation Systems

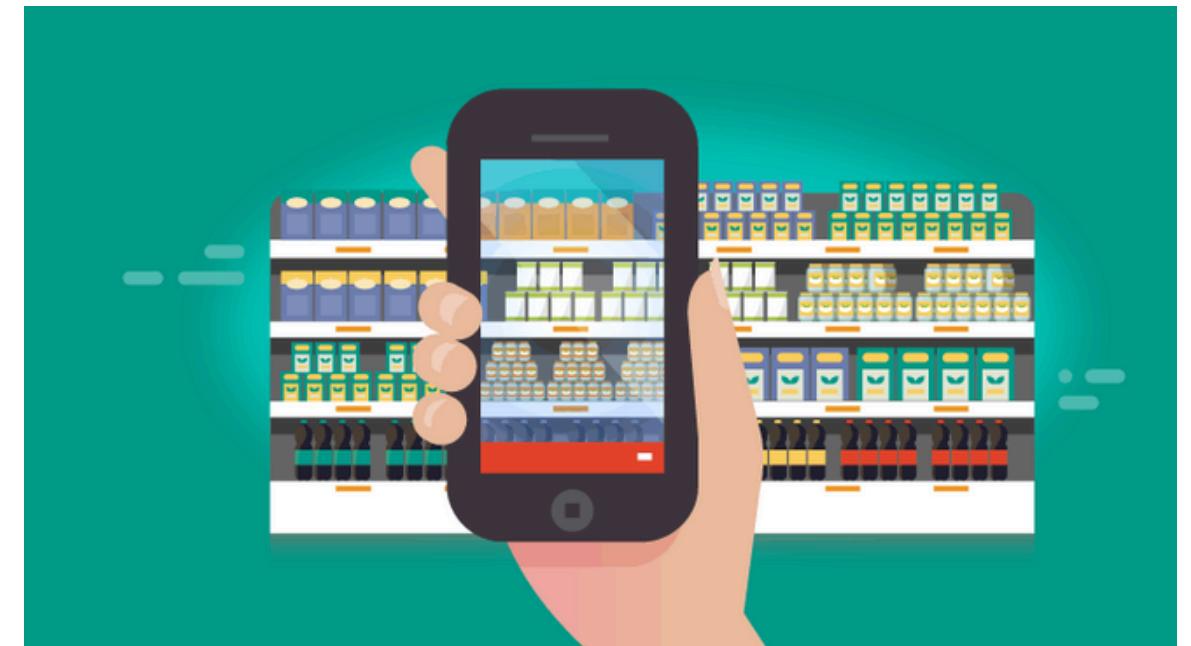
COLLABORATIVE FILTERING



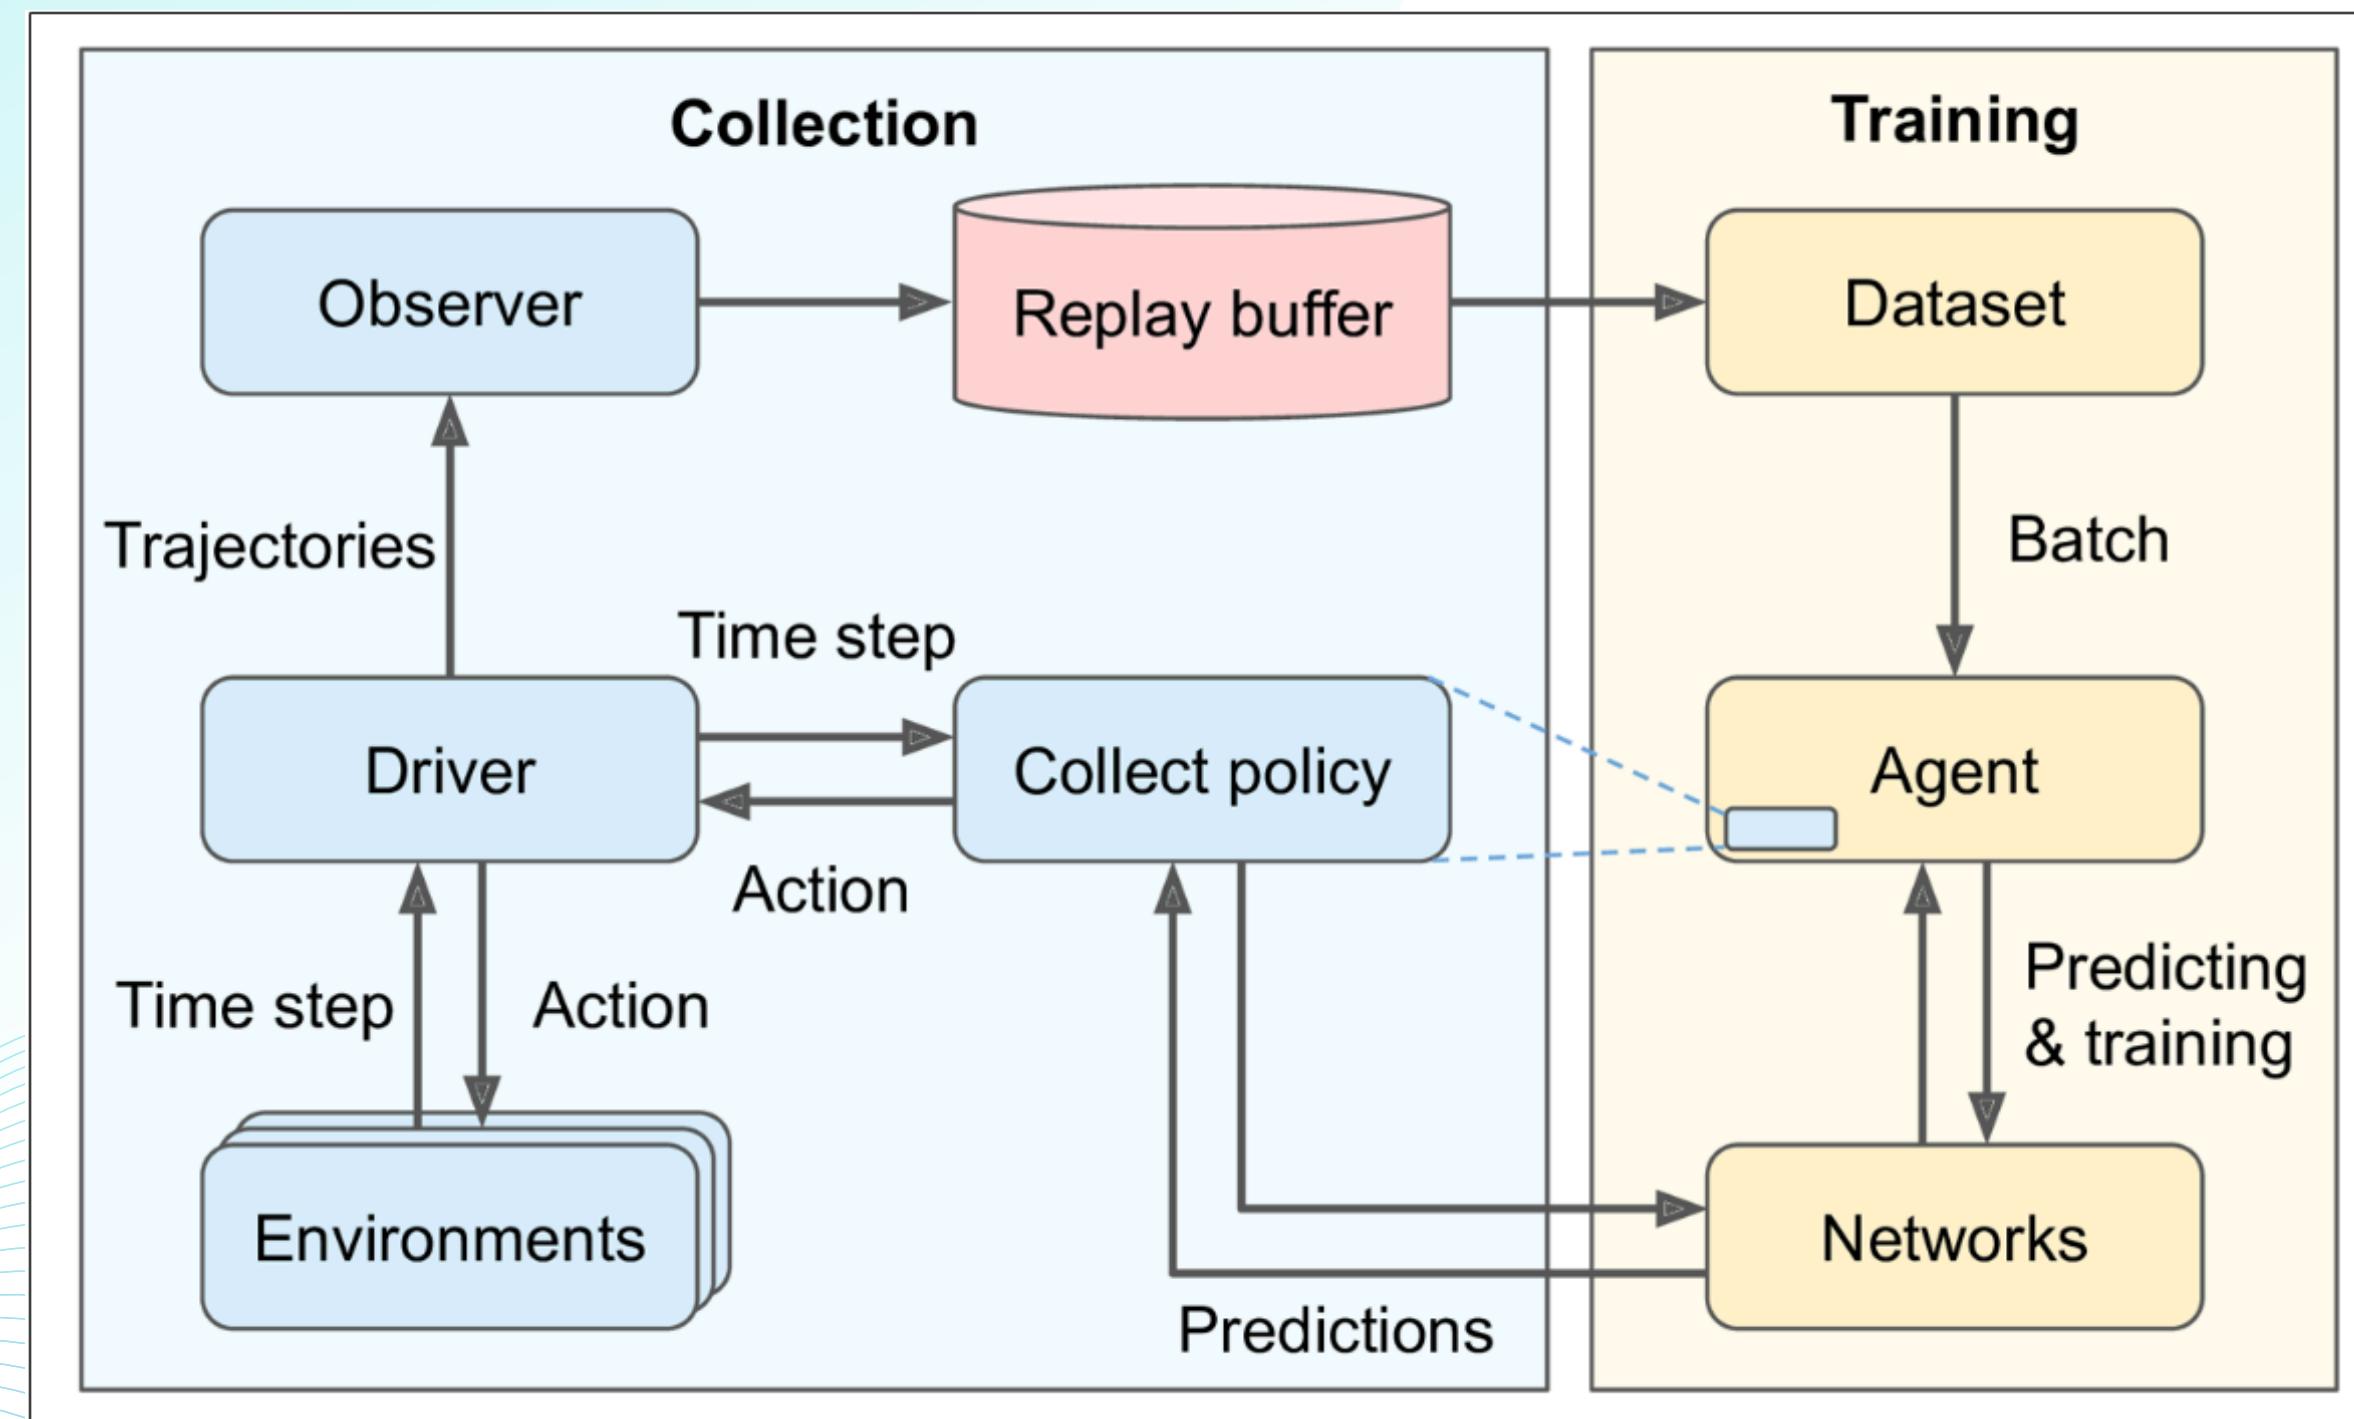
Medical Diagnostics



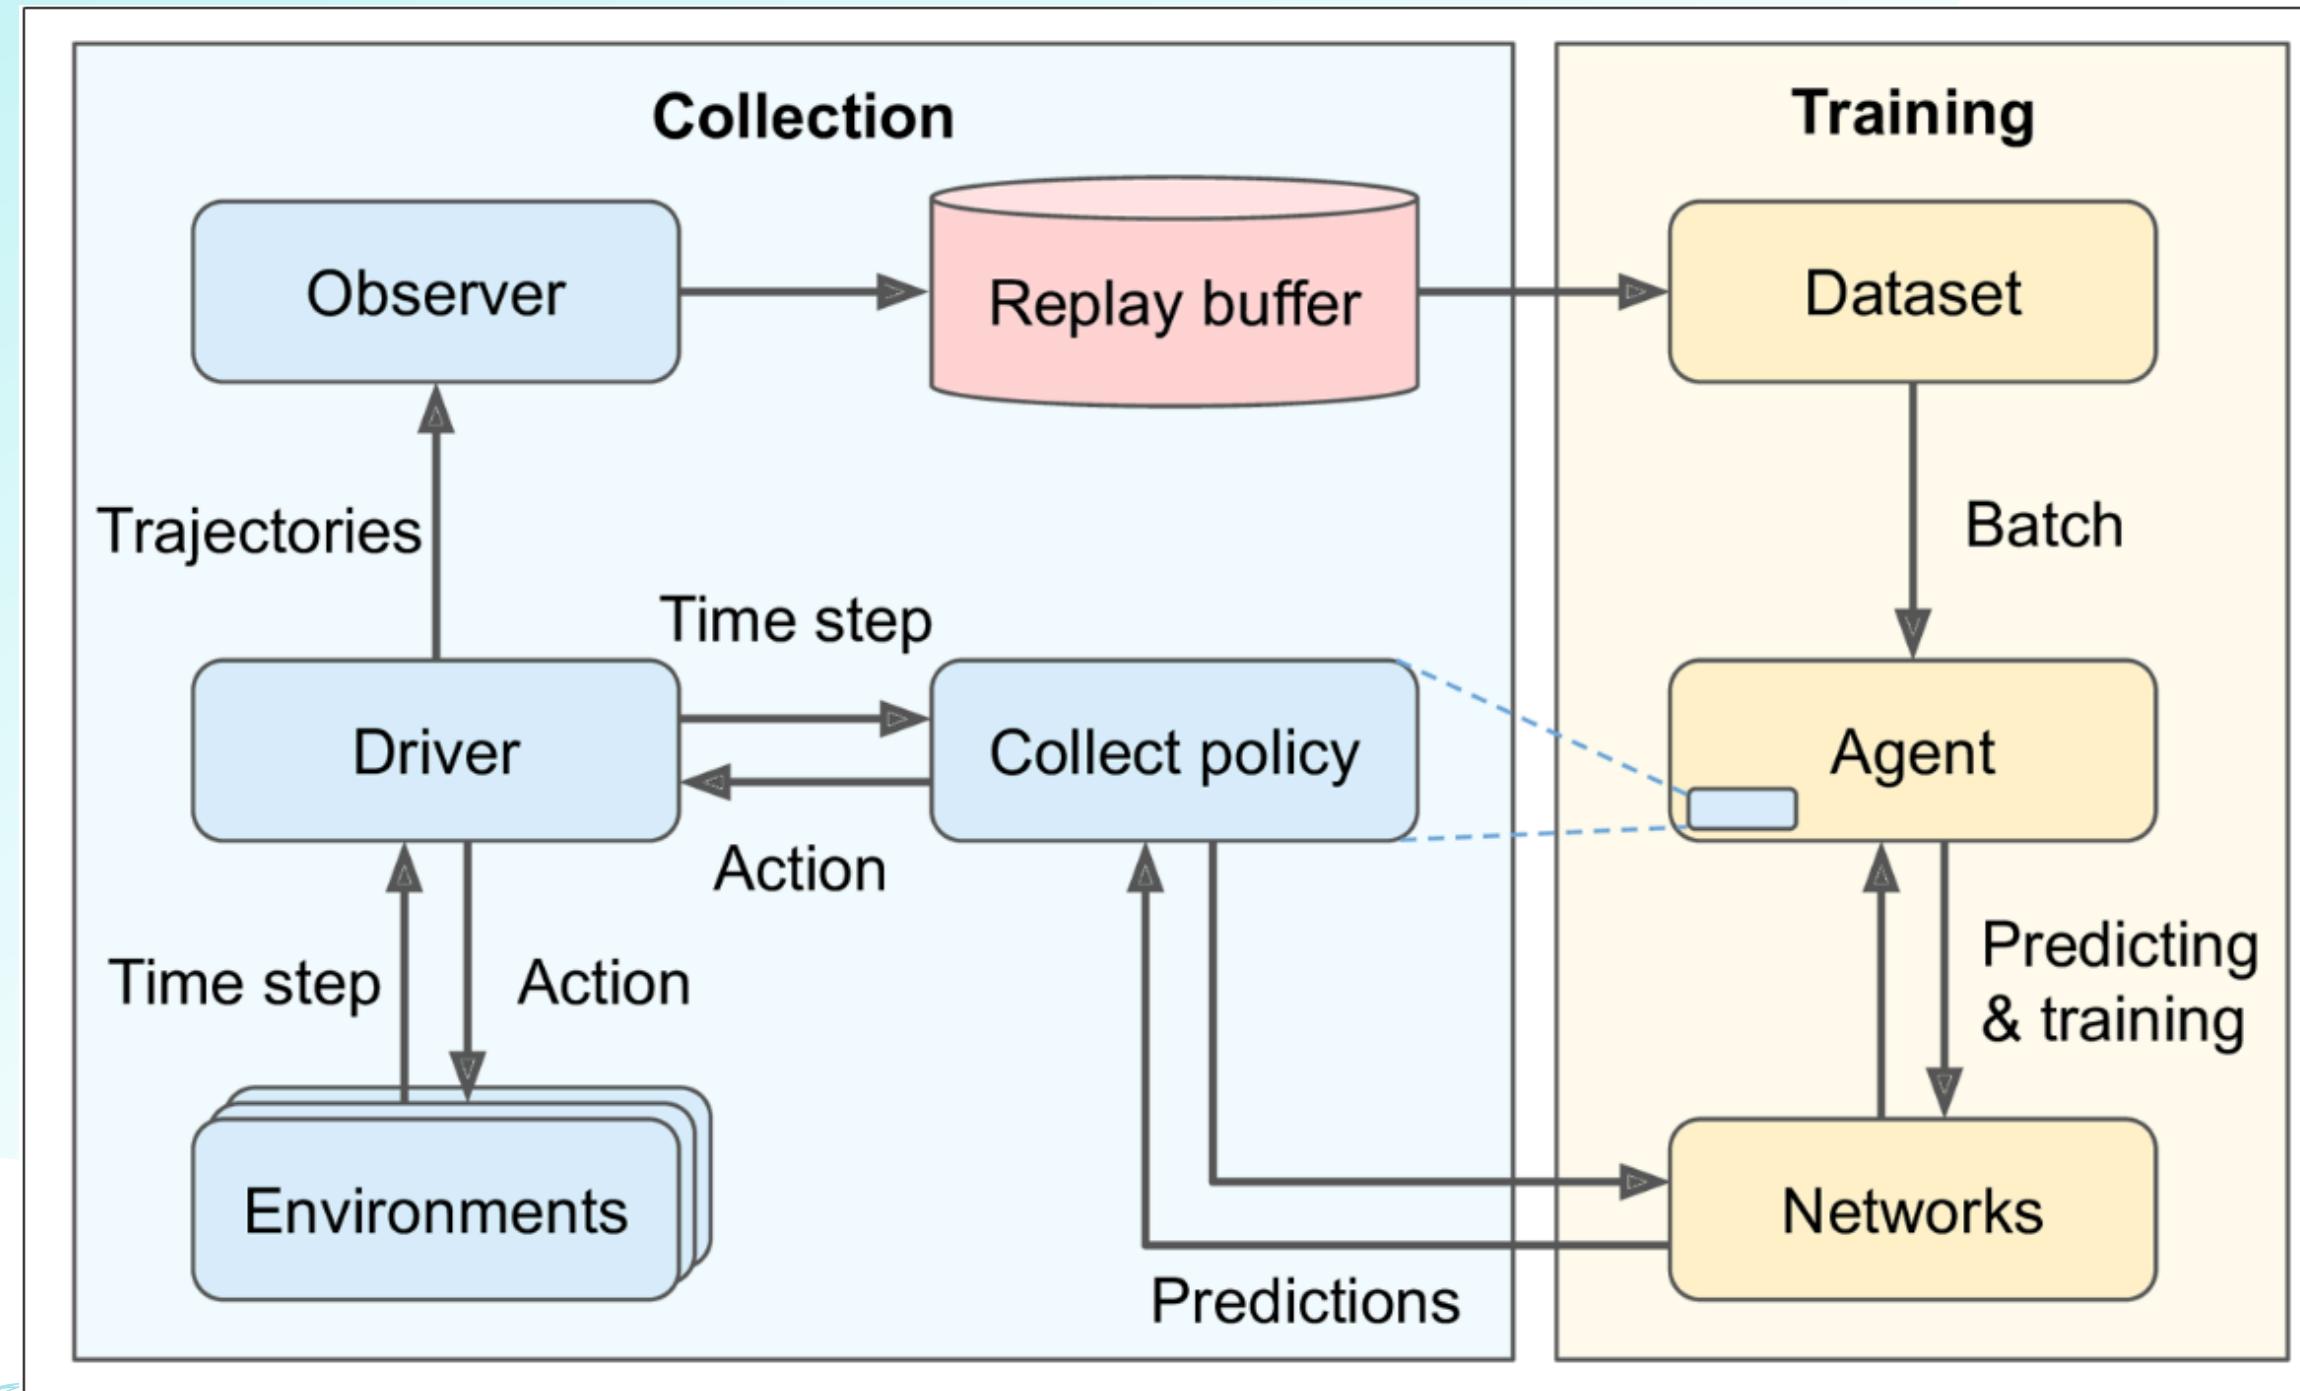
Image Recognition



ARQUITECTURE



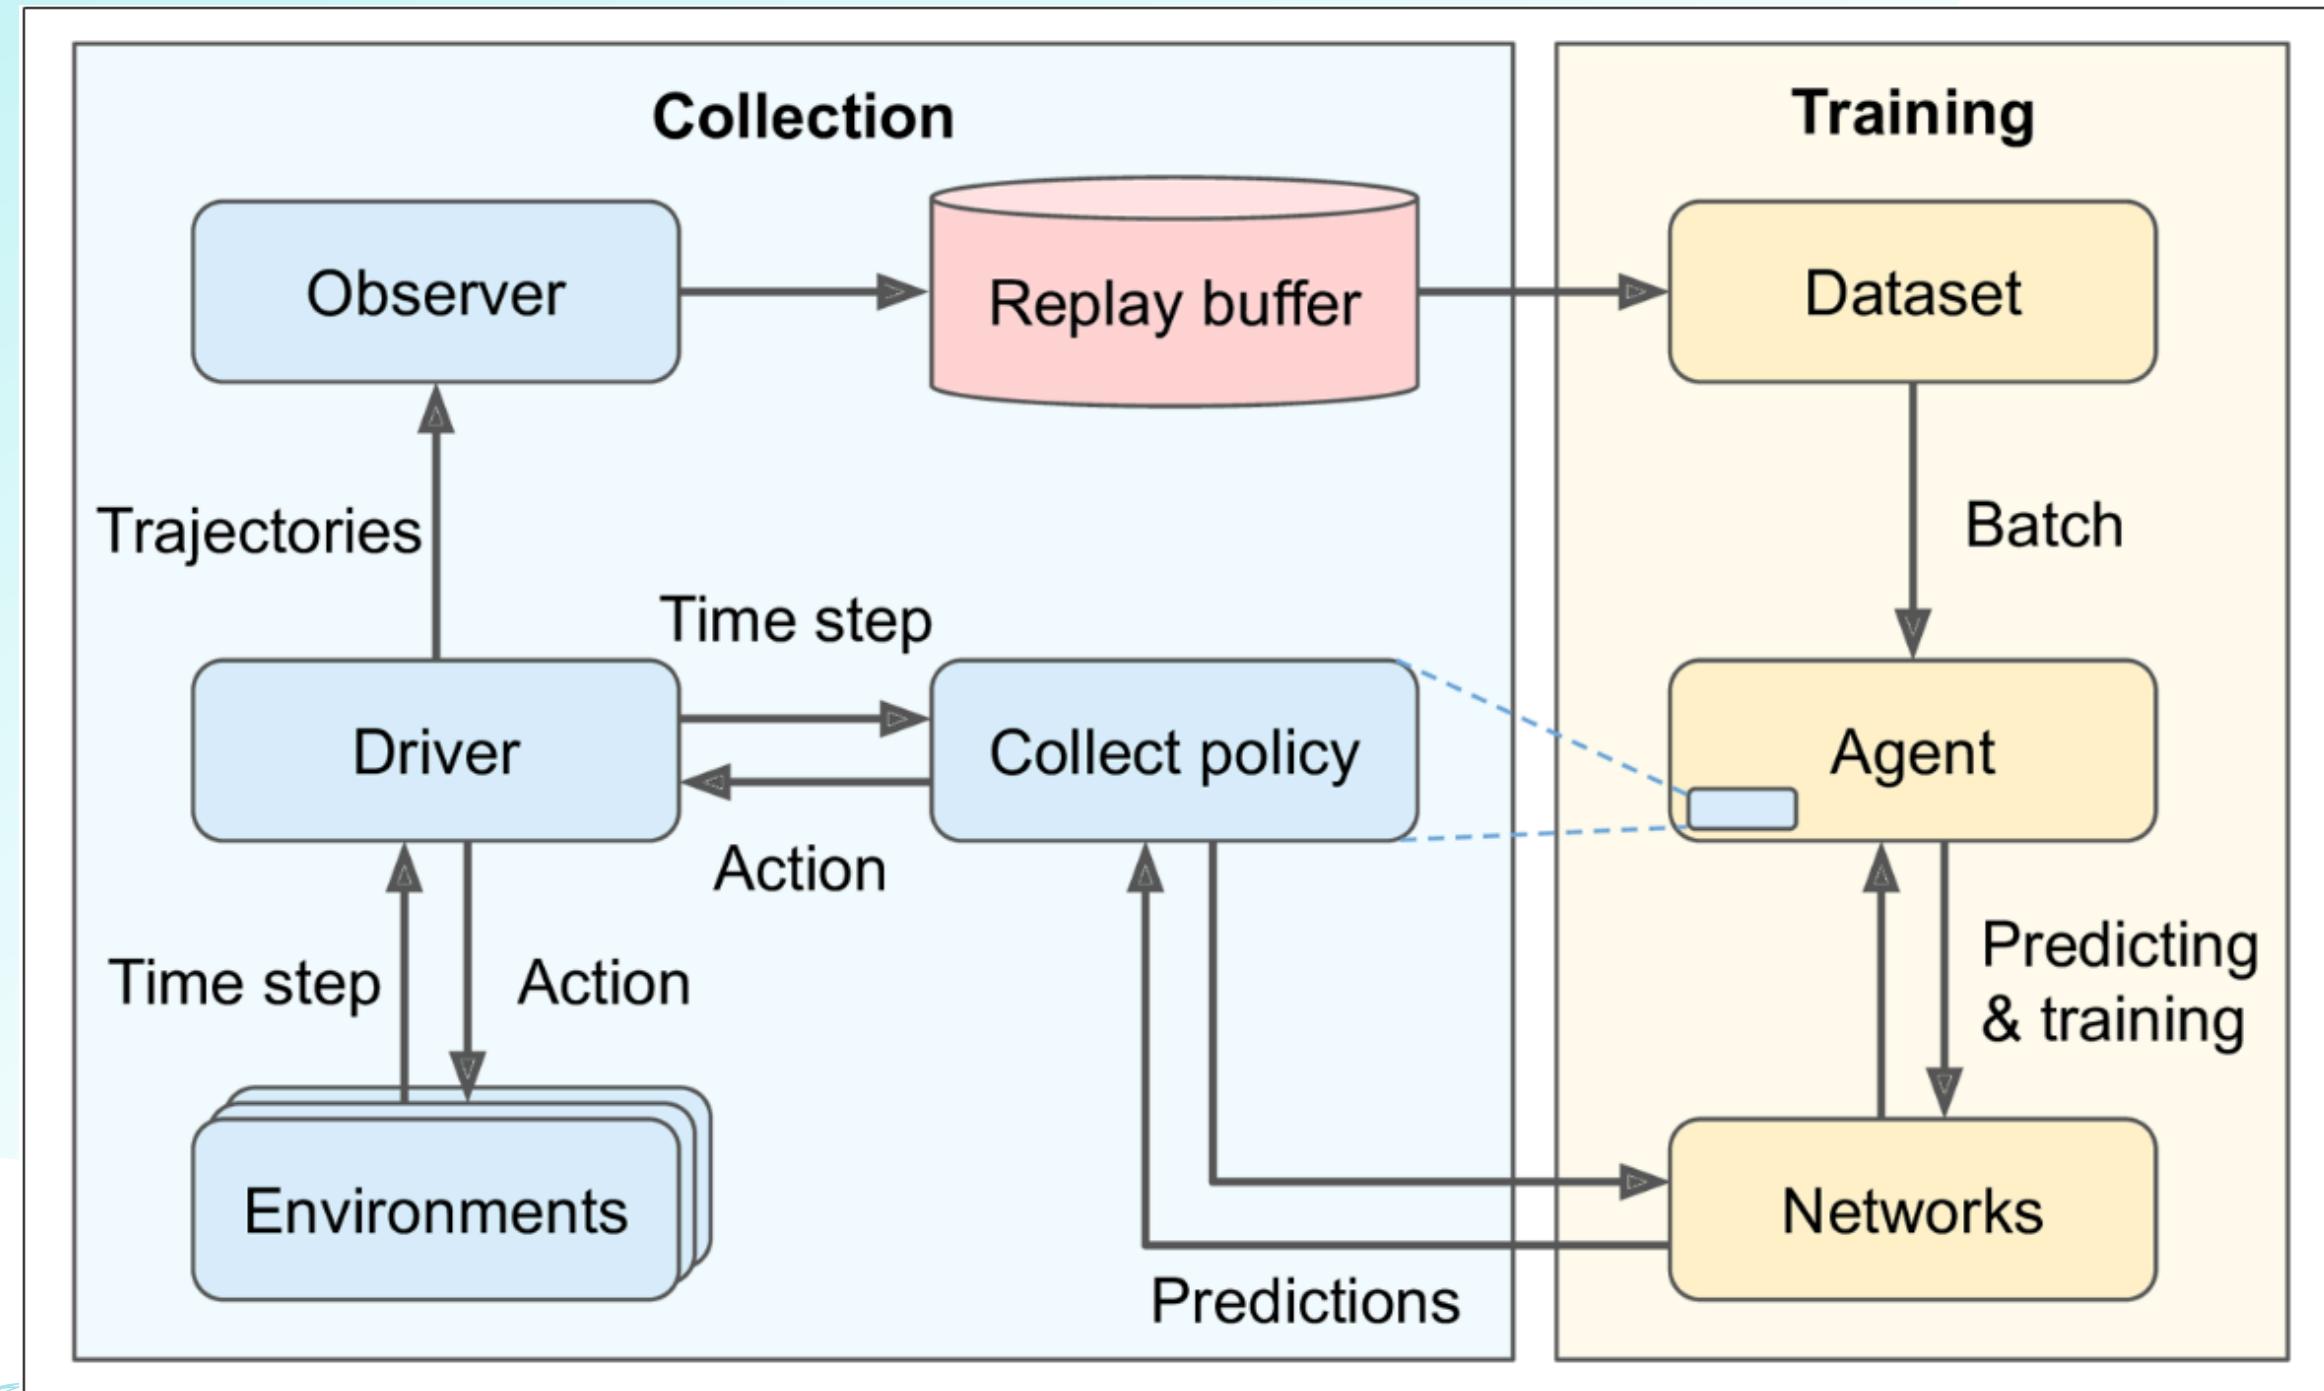
ARQUITECTURE



Collection

- **Environments:**
 - Provides time step after each action;
- **Driver:**
 - Executes actions and receives feedback;
- **Collection Policy:**
 - Used policy during data collection;
- **Observer:**
 - Collects Trajectories;
- **Replay Buffer:**
 - Stores Trajectories.

ARQUITECTURE

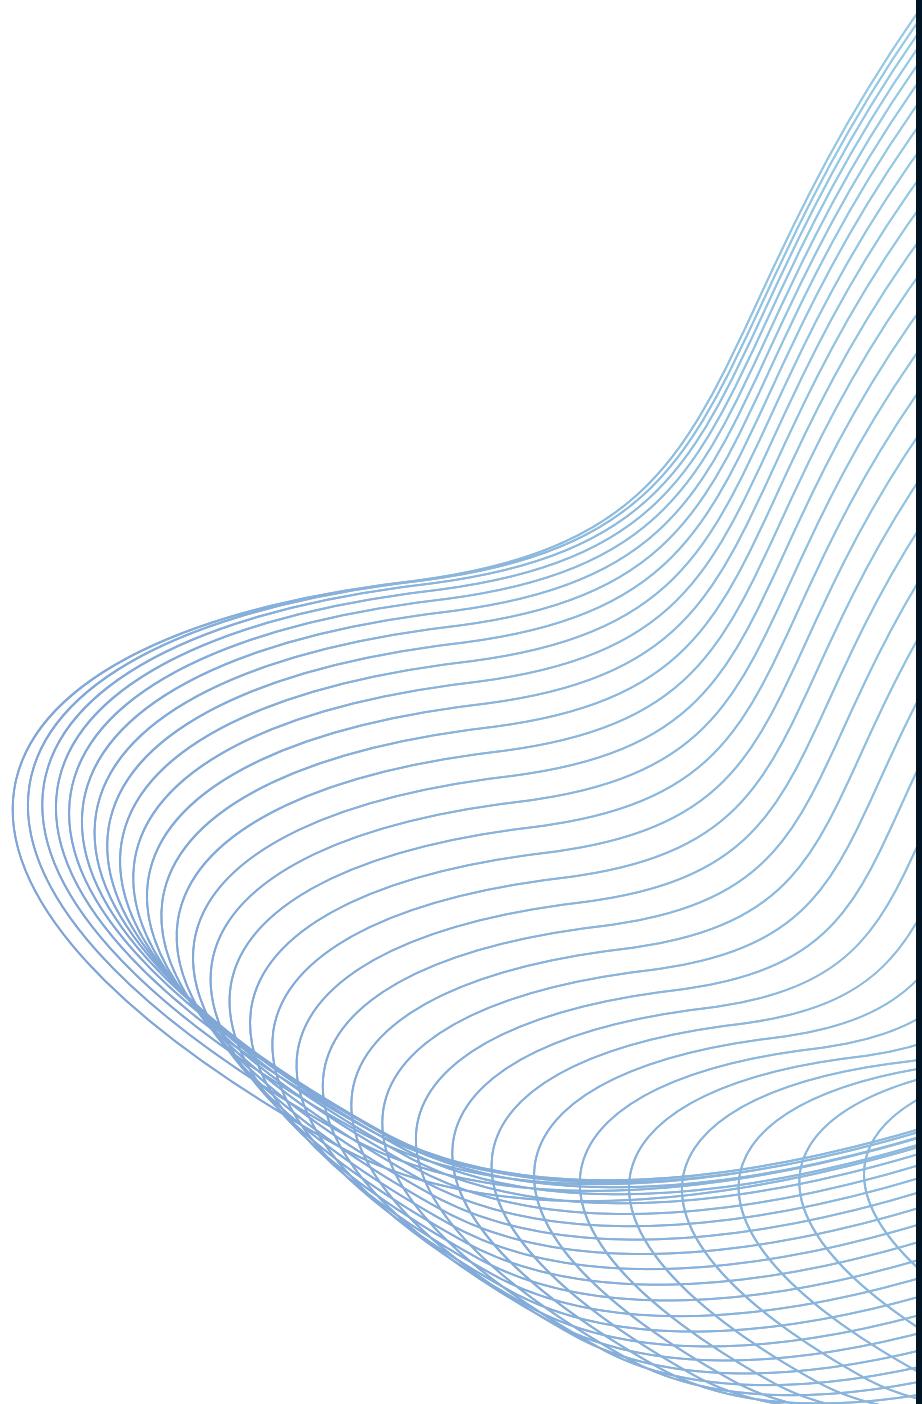


Training

- **Dataset**
 - Generates Batches;
- **Agent:**
 - Uses collected data to adjust the policy and networks;
 - Supports several RL algorithms;
- **Networks:**
 - Represent the models used for predictions.

ENVIRONMENTS

- **Observation:** This is the part of the environment state that the agent can observe to choose its actions at the next step.
- **Reward:** The agent is learning to maximize the sum of these rewards across multiple steps
- **step_type:** The first, intermediate or last time step in a sequence
- **discount:** how much to weight the reward at the next time step relative to the reward at the current time step



```
import abc

class PyEnvironment(object):

    def reset(self):
        """Return initial_time_step."""
        self._current_time_step = self._reset()
        return self._current_time_step

    def step(self, action):
        """Apply action and return new_time_step."""
        if self._current_time_step is None:
            return self.reset()
        self._current_time_step = self._step(action)
        return self._current_time_step

    def current_time_step(self):
        return self._current_time_step

    def time_step_spec(self):
        """Return time_step_spec."""

    @abc.abstractmethod
    def observation_spec(self):
        """Return observation_spec."""

    @abc.abstractmethod
    def action_spec(self):
        """Return action_spec."""

    @abc.abstractmethod
    def _reset(self):
        """Return initial_time_step."""

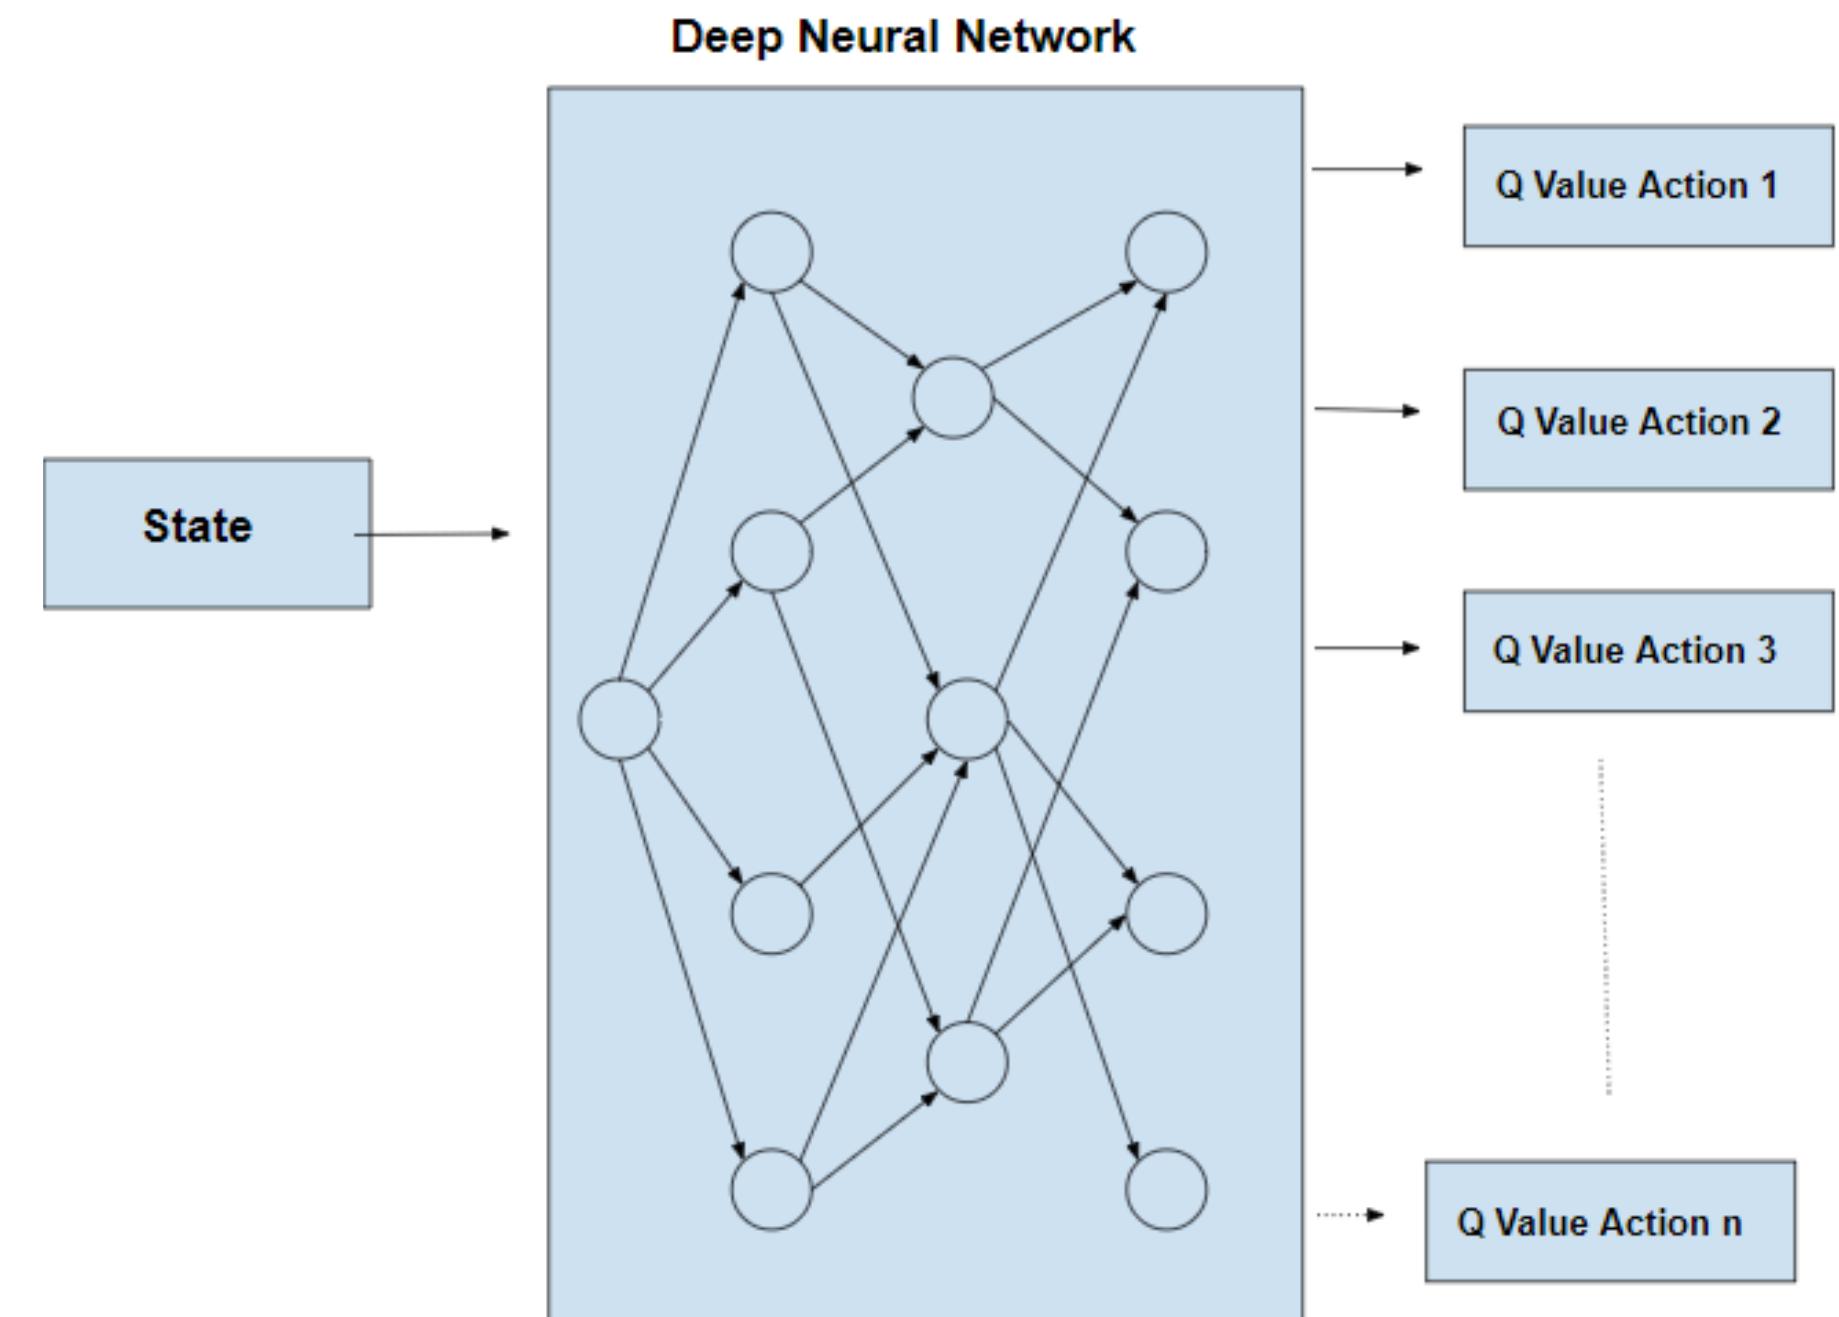
    @abc.abstractmethod
    def _step(self, action):
        """Apply action and return new_time_step."""
```

NEURAL NETWORKS

TensorFlow Main Networks : QNetwork, CriticNetworks, ActorNetworks, ActorDistributionNetworks

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

discount factor = 0.90
immediate reward future reward



SETTING UP THE DQN AGENT

```
# 1. Configure the environment
env_name = "CartPole-v1" # Environment name
train_env = TFPyEnvironment(suite_gym.load(env_name)) # Training environment
eval_env = suite_gym.load(env_name) # Evaluation environment (non-TF)

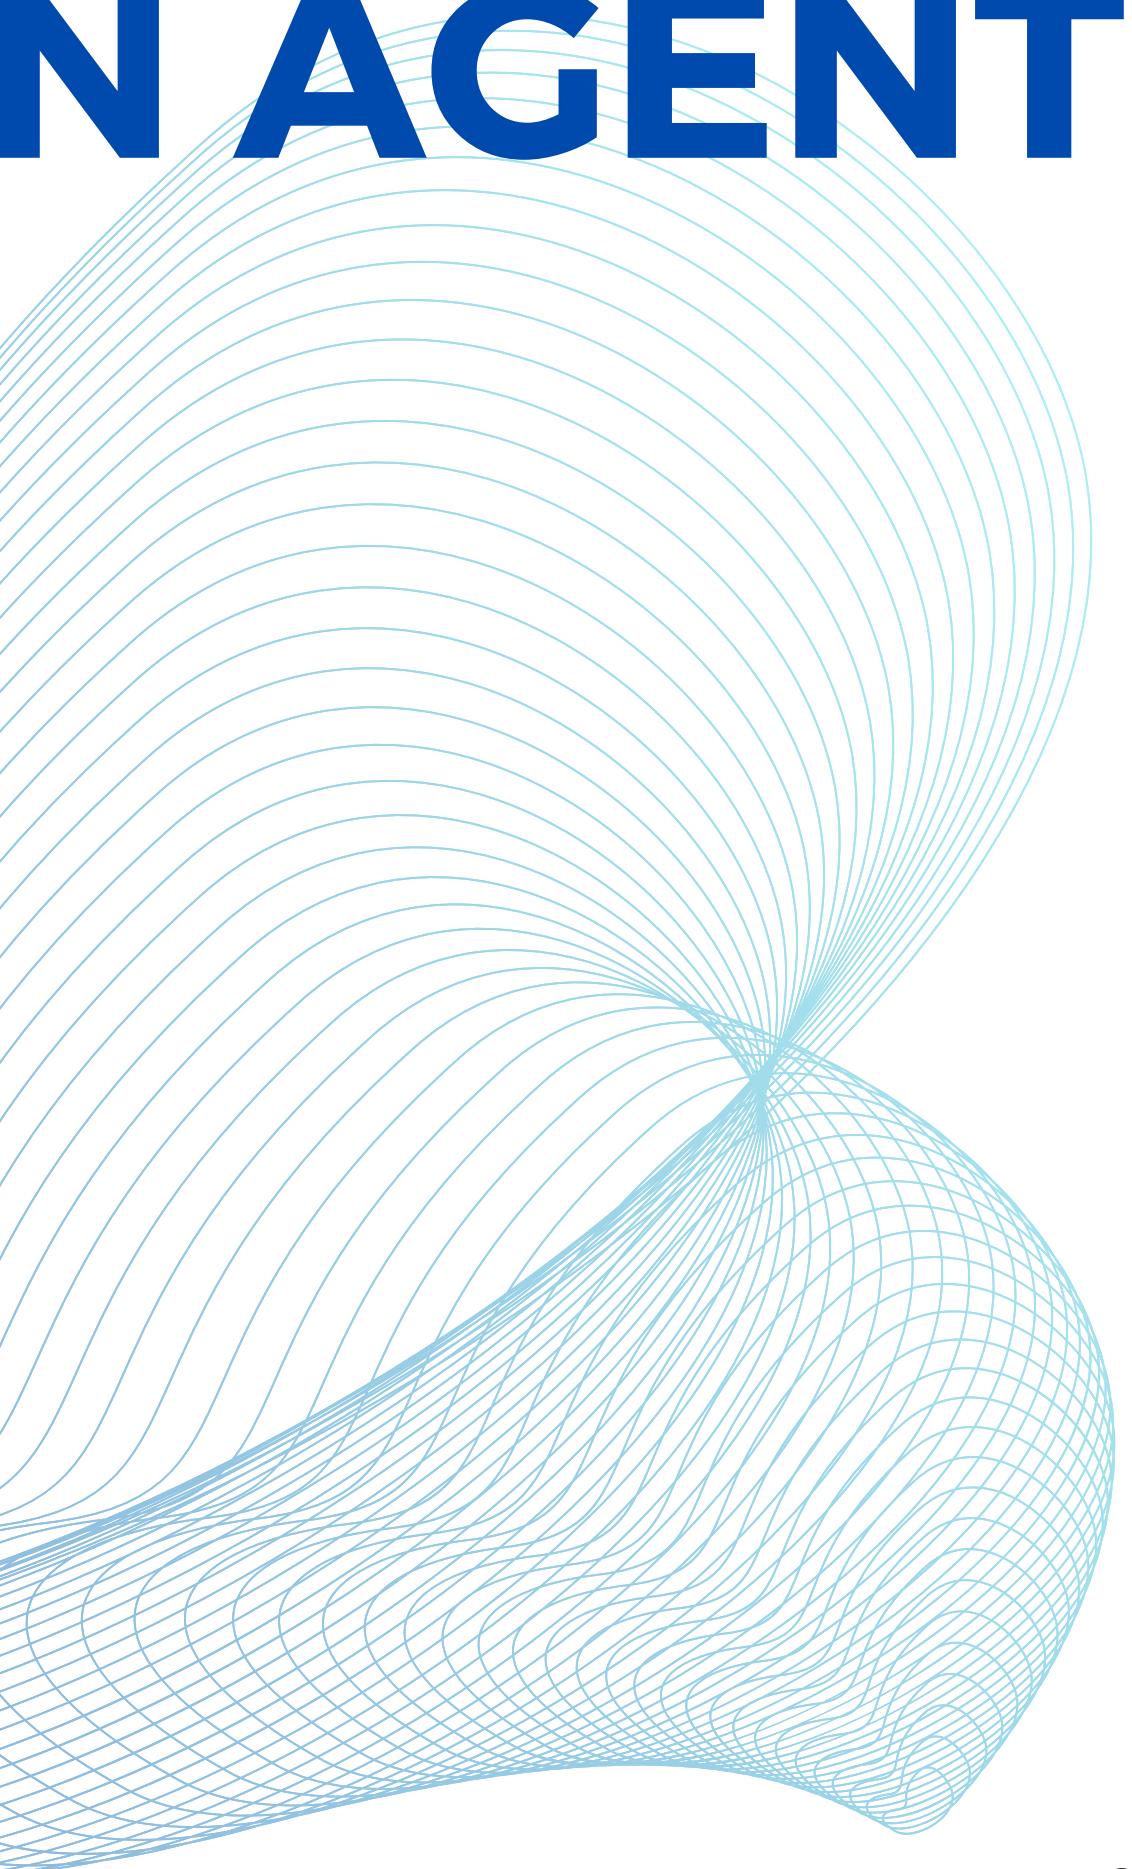
# 2. Create the Q network and agent
fc_layer_params = (100,) # Neural network parameters

q_net = QNetwork(
    train_env.observation_spec(),
    train_env.action_spec(),
    fc_layer_params=fc_layer_params
)

optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)

agent = DqnAgent(
    train_env.time_step_spec(),
    train_env.action_spec(),
    q_network=q_net,
    optimizer=optimizer,
    td_errors_loss_fn=tf.keras.losses.Huber(reduction="none"),
    train_step_counter=tf.Variable(0)
)
agent.initialize()

# 3. Configure the replay buffer and data collection
replay_buffer = TFUniformReplayBuffer(
    data_spec=agent.collect_data_spec,
    batch_size=train_env.batch_size,
    max_length=10000
)
```



METRICS

Plot Graphs:



TensorBoard

```
# Define the metrics + TensorBoard writer
train_metrics = [
    tf.metrics.NumberOfEpisodes(),
    tf.metrics.EnvironmentSteps(),
    tf.metrics.AverageReturnMetric(),
    tf.metrics.AverageEpisodeLengthMetric(),
]
train_summary_writer = tf.summary.create_file_writer(log_dir)
train_summary_writer.set_as_default()
```

DEMO

THANKS FOR LISTENING

Do you have any question?

